

Lecture 4

Orthogonal Projections; Regression and Linear Algebra

DSC 40A, Summer 2026

Announcements at break

Lecture 4 Part 1

Orthogonal Projections

DSC 40A, Summer 2026

Agenda for Part 1

- Spans and projections.
- Matrices.
- Spans and projections, revisited.
- Regression and linear algebra.

Question 🤔

Answer at q.dsc40a.com

Remember, you can always ask questions at q.dsc40a.com!

If the direct link doesn't work, click the "🤔 Lecture Questions"
link in the top right corner of dsc40a.com.

Spans and projections

Recap


Projecting onto a single vector

- Let $\vec{x}$ and $\vec{y}$ be two vectors in $\mathbb{R}^n$.
- The span of $\vec{x}$ is the set of all vectors of the form:

$$w\vec{x}$$

where $w \in \mathbb{R}$ is a scalar.


- Question:** What vector in $\text{span}(\vec{x})$ is closest to $\vec{y}$?
- The vector in $\text{span}(\vec{x})$ that is closest to $\vec{y}$ is the **projection of $\vec{y}$ onto $\text{span}(\vec{x})$** .



Projection error

Find minimum
Small as possible $\rightarrow$ minimizing MSE




- Let $\vec{e} = \vec{y} - w\vec{x}$ be the **projection error**:
that is, the vector that connects $\vec{y}$ to $\text{span}(\vec{x})$.
- Goal: Find the w that makes $\vec{e}$ as short as possible.**
 - That is, minimize:
 $\|\vec{e}\|$
 - Equivalently, minimize:
 $\|\vec{y} - w\vec{x}\|$
- Idea:** To make $\vec{e}$ as short as possible, it should be **orthogonal** to $w\vec{x}$.



$$\Rightarrow x = y + 2 \Rightarrow$$

Minimizing projection error

- Goal: Find the w that makes $\vec{e} = \vec{y} - w\vec{x}$ as short as possible.
- Idea: To make $\vec{e}$ as short as possible, it should be orthogonal to $w\vec{x}$.
- Can we prove that making $\vec{e}$ orthogonal to $w\vec{x}$ minimizes $\|\vec{e}\|$?



$$\|\vec{e}\|^2 = \|\vec{e}^*\|^2 + b^2$$

best is shortest + something

$$\|\vec{e}\|^2 = \|\vec{e}^*\|^2 + b^2$$

$$\|\vec{e}^*\|^2 \leq \|\vec{e}\|^2$$

$$\|\vec{e}\|^2 \geq \|\vec{e}^*\|^2$$

$e^* \leq e^n + e^*$ orthogonal to projection: smallest $\vec{e}$ is orthogonal

$w\vec{x} \perp \vec{e}$ $w\vec{x} \perp \vec{x}$ are in the same direction
Scalar

$w\vec{x} \perp \vec{e}$ = same angle

Minimizing projection error

- Goal: Find the w that makes $\vec{e} = \vec{y} - w\vec{x}$ as short as possible.
- Now we know that to minimize $\|\vec{e}\|$, $\vec{e}$ must be orthogonal to $w\vec{x}$.
- Given this fact, how can we solve for w ?

$\vec{e}$ orthogonal to $w\vec{x} \Rightarrow w\vec{x} \cdot \vec{e} = 0$

$$w\vec{x} \cdot (\vec{y} - w\vec{x}) = 0$$


$$\vec{x} \cdot (\vec{y} - w\vec{x}) = 0$$

$$\vec{x} \cdot \vec{y} - \vec{x} \cdot (w\vec{x}) = 0$$

$$\vec{x} \cdot \vec{y} - w(\vec{x} \cdot \vec{x}) = 0$$

$$w = \frac{\vec{x} \cdot \vec{y}}{\vec{x} \cdot \vec{x}}$$

$$w^* = \frac{\vec{x} \cdot \vec{y}}{\vec{x} \cdot \vec{x}}$$



Orthogonal projection

Goal

• Question: What vector in $\text{span}(\vec{x})$ is closest to $\vec{y}$?

• Answer: It is the vector $w^*\vec{x}$, where:

$$w^* = \frac{\vec{x} \cdot \vec{y}}{\vec{x} \cdot \vec{x}} \quad \nearrow \text{MSE}$$


• Note that w^* is the solution to a minimization problem, specifically, this one:

verbal

$$\text{error}(w) = \|\vec{e}\| = \|\vec{y} - w\vec{x}\|$$

• We call $w^*\vec{x}$ the **orthogonal projection** of $\vec{y}$ onto $\text{span}(\vec{x})$.

◦ Think of $w^*\vec{x}$ as the "shadow" of $\vec{y}$.



perpendicular
in ml. dim.

Exercise

Let $\vec{a} = \begin{bmatrix} 5 \\ 2 \end{bmatrix}$ and $\vec{b} = \begin{bmatrix} -1 \\ 9 \end{bmatrix}$.


Vector in
 $\text{span}(\vec{b})$ closest to
 $\vec{a}$

linear
combination
of $\vec{b}$

What is the orthogonal projection of $\vec{a}$ onto $\text{span}(\vec{b})$?

Your answer should be of the form $w^* \vec{b}$, where w^* is a scalar.

$$\begin{aligned} w^* &= \frac{\vec{a} \cdot \vec{b}}{\vec{b} \cdot \vec{b}} = \frac{\vec{b} \cdot \vec{a}}{\vec{b} \cdot \vec{b}} \quad \text{projecting onto} \\ &= \frac{(-1)(5) + (9)(2)}{(-1)(-1) + (9)^2} = \frac{13}{82} \end{aligned}$$



Orthogonal projection
of $\vec{a}$ onto $\text{span}(\vec{b})$
is $w^* \vec{b} = \frac{13}{82} \vec{b}$


$\text{Span}(\vec{b})$ line

$\text{Span}(\vec{b}, \vec{a})$

Moving to multiple dimensions

- Let's now consider three vectors, $\vec{y}$, $\vec{x}^{(1)}$, and $\vec{x}^{(2)}$, all in $\mathbb{R}^n$.
- Question:** What vector in $\text{span}(\vec{x}^{(1)}, \vec{x}^{(2)})$ is closest to $\vec{y}$?
 - Vectors in $\text{span}(\vec{x}^{(1)}, \vec{x}^{(2)})$ are of the form $w_1 \vec{x}^{(1)} + w_2 \vec{x}^{(2)}$, where $w_1, w_2 \in \mathbb{R}$ are scalars.
- Before trying to answer, let's watch 🎥 [this animation that Jack, one of our tutors, made](#).

projection
 wx



Minimizing projection error in multiple dimensions

span a plane

- **Question:** What vector in $\text{span}(\vec{x}^{(1)}, \vec{x}^{(2)})$ is closest to $\vec{y}$?

- That is, what vector minimizes $\|\vec{e}\|$, where:

$$\vec{e} = \vec{y} - w_1 \vec{x}^{(1)} - w_2 \vec{x}^{(2)}$$

- **Answer:** It's the vector such that $w_1 \vec{x}^{(1)} + w_2 \vec{x}^{(2)}$ is orthogonal to $\vec{e}$.

- **Issue:** Solving for w_1 and w_2 in the following equation is difficult:

$$(w_1 \vec{x}^{(1)} + w_2 \vec{x}^{(2)}) \cdot \underbrace{(\vec{y} - w_1 \vec{x}^{(1)} - w_2 \vec{x}^{(2)})}_{\vec{e}} = 0$$

any
vector

error
of

Minimizing projection error in multiple dimensions

- It's hard for us to solve for w_1 and w_2 in:

$$\left(w_1 \vec{x}^{(1)} + w_2 \vec{x}^{(2)} \right) \cdot \underbrace{\left(\vec{y} - w_1 \vec{x}^{(1)} - w_2 \vec{x}^{(2)} \right)}_{\vec{e}} = 0$$

- Observation:** All we really need is for $\vec{x}^{(1)}$ and $\vec{x}^{(2)}$ to individually be orthogonal to $\vec{e}$.

- That is, it's sufficient for $\vec{e}$ to be orthogonal to the spanning vectors themselves.

- If $\vec{x}^{(1)} \cdot \vec{e} = 0$ and $\vec{x}^{(2)} \cdot \vec{e} = 0$, then:

$$\begin{aligned} (w_1 \vec{x}^{(1)} + w_2 \vec{x}^{(2)}) \cdot \vec{e} &= w_1 \vec{x}^{(1)} \cdot \vec{e} + w_2 \vec{x}^{(2)} \cdot \vec{e} \\ &= w_1 (\vec{x}^{(1)} \cdot \vec{e}) + w_2 (\vec{x}^{(2)} \cdot \vec{e}) \\ &= w_1 (0) + w_2 (0) = 0 \end{aligned}$$

Minimizing projection error in multiple dimensions

- **Question:** What vector in $\text{span}(\vec{x}^{(1)}, \vec{x}^{(2)})$ is closest to $\vec{y}$?
- **Answer:** It's the vector such that $w_1\vec{x}^{(1)} + w_2\vec{x}^{(2)}$ is **orthogonal** to $\vec{e} = \vec{y} - w_1\vec{x}^{(1)} - w_2\vec{x}^{(2)}$.
- Equivalently, it's the vector such that $\vec{x}^{(1)}$ and $\vec{x}^{(2)}$ are both orthogonal to $\vec{e}$:

$$\begin{array}{l} \vec{x}^{(1)} \cdot \vec{e} \\ \vec{x}^{(2)} \cdot \vec{e} \end{array} = \begin{array}{l} \vec{x}^{(1)} \cdot \left(\vec{y} - w_1\vec{x}^{(1)} - w_2\vec{x}^{(2)} \right) = 0 \\ \vec{x}^{(2)} \cdot \left(\vec{y} - w_1\vec{x}^{(1)} - w_2\vec{x}^{(2)} \right) = 0 \end{array}$$

$\underbrace{\vec{y} - w_1\vec{x}^{(1)} - w_2\vec{x}^{(2)}}_{\vec{e}}$

w_1^*
 w_2^*

- This is a system of two equations, two unknowns (w_1 and w_2), but it still looks difficult to solve.

Now what?

- We're looking for the scalars w_1 and w_2 that satisfy the following equations:

$$\begin{aligned}\vec{x}^{(1)} \cdot \left(\vec{y} - w_1 \vec{x}^{(1)} - w_2 \vec{x}^{(2)} \right) &= 0 \\ \vec{x}^{(2)} \cdot \left(\vec{y} - w_1 \vec{x}^{(1)} - w_2 \vec{x}^{(2)} \right) &= 0\end{aligned}$$

Handwritten notes: $\vec{x}^{(1)}$ and $\vec{x}^{(2)}$ are written in blue below the first equation. A bracket under the second equation points to a red $\vec{e}$ below it.

- In this example, we just have two spanning vectors, $\vec{x}^{(1)}$ and $\vec{x}^{(2)}$.
- If we had any more, this system of equations would get extremely messy, extremely quickly.

 **Idea:** Rewrite the above system of equations as a single equation, involving matrix-vector products.

Matrices

Matrices

$\mathbb{R}^2$ all real #s

- An $n \times d$ **matrix** is a table of numbers with n rows and d columns.
- We use upper-case letters to denote matrices.

$\rightarrow A = \begin{bmatrix} 2 & 5 & 8 \\ -1 & 5 & -3 \end{bmatrix}$

- Since A has two rows and three columns, we say $A \in \mathbb{R}^{2 \times 3}$.
- **Key idea:** Think of a matrix as **several column vectors, stacked next to each other**.

set of matrices
w/ 2 rows
& 3 cols

$$A = \begin{bmatrix} 2 \\ -1 \end{bmatrix} \begin{bmatrix} 5 \\ 5 \end{bmatrix} \begin{bmatrix} 8 \\ -3 \end{bmatrix}$$

Matrix addition and scalar multiplication

- We can add two matrices only if they have the same dimensions.

- Addition occurs elementwise:

$$\begin{bmatrix} 2 & 5 & 8 \\ -1 & 5 & -3 \end{bmatrix} + \begin{bmatrix} 1 & 2 & 3 \\ 0 & 1 & 2 \end{bmatrix} = \begin{bmatrix} 3 & 7 & 11 \\ -1 & 6 & -1 \end{bmatrix}$$

*same # of rows
same # of cols*

- Scalar multiplication occurs elementwise, too:

$$2 \begin{bmatrix} 2 & 5 & 8 \\ -1 & 5 & -3 \end{bmatrix} = \begin{bmatrix} 4 & 10 & 16 \\ -2 & 10 & -6 \end{bmatrix}$$

$$\vec{V} = \begin{pmatrix} 1 \\ 2 \end{pmatrix}_{2 \times 1}, \quad \vec{V}^T = \begin{pmatrix} 1 & 2 \end{pmatrix}_{1 \times 2}$$

Matrix-matrix multiplication

- **Key idea:** We can multiply matrices A and B if and only if:

$\# \text{ columns in } A = \# \text{ rows in } B$

- If A is $n \times d$ and B is $d \times p$, then AB is $n \times p$.

- Example: If A is as defined below, what is $A^T A$?

$$A^T = \begin{bmatrix} 1 & 2 \\ 2 & 5 \\ 8 & -3 \end{bmatrix}_{3 \times 2}$$

(Note: The first row of A^T is circled in pink in the original image, corresponding to the first row of A .)

$$A = \begin{bmatrix} 2 & 5 & 8 \\ -1 & 5 & -3 \end{bmatrix}_{2 \times 3}$$

(Note: The first column of A is circled in pink in the original image, corresponding to the first row of A^T .)

$$A^T A = \begin{bmatrix} 5 & - & - \\ - & - & - \\ - & - & 25 \end{bmatrix}_{3 \times 3}$$

(Note: The top-left element 5 and the bottom-right element 25 are boxed in yellow and blue respectively in the original image.)

$$\begin{pmatrix} 2 & -1 \end{pmatrix} \cdot \begin{pmatrix} 2 \\ -1 \end{pmatrix}$$

$$2 \cdot 2 + (-1) \cdot (-1) = 5$$

Question 🤔

Answer in chat

Assume A , B , and C are all matrices. Select the **incorrect** statement below.

• A. $A(B + C) = AB + AC$.

• B. $A(BC) = (AB)C$.

• C. $AB = BA$.

• D. $(A + B)^T = A^T + B^T$.

• E. $(AB)^T = B^T A^T$.

2×3 3×7
 $A \neq B$

$\dim(A) \neq \dim(B)$

AB
 2×3 3×7

BA
 3×7 2×3

$A_{3 \times 7} B_{7 \times 5} = 3 \times 5$
 $B_{7 \times 5} A_{5 \times 7} = 7 \times 7$

Matrix-vector multiplication

- A vector $\vec{v} \in \mathbb{R}^n$ is a matrix with n rows and 1 column.

$$\vec{v} = \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{bmatrix} \quad n \times 1$$

- Suppose $A \in \mathbb{R}^{n \times d}$.

- What must the dimensions of $\vec{v}$ be in order for the product $A\vec{v}$ to be valid?

$A_{n \times d} \leftarrow \vec{v}_{d \times 1} \quad \vec{v} \in \mathbb{R}^d$ d components to be vector right

- What must the dimensions of $\vec{v}$ be in order for the product $\vec{v}^T A$ to be valid?

$\left[\begin{array}{ccc} \vdots & \dots & \vdots \end{array} \right] \left[\begin{array}{c} \vdots \\ \vdots \\ \vdots \end{array} \right] \quad V_{1 \times n}^T \quad A_{n \times d} \quad \vec{v} \in \mathbb{R}^n \quad \vec{v} = \begin{bmatrix} 1 \\ \vdots \\ n \end{bmatrix} \quad \vec{v}^T = \begin{bmatrix} 1 & \dots & n \end{bmatrix}$

One view of matrix-vector multiplication

- One way of thinking about the product $A\vec{v}$ is that it is the **dot product of $\vec{v}$ with every row of A** .
- Example: What is $A\vec{v}$?

Handwritten annotations: 2×3 (yellow), 3×1 (yellow), $\mathbb{R}^3$ (red), $A\vec{v} \in \mathbb{R}^2$ (red).

$$A = \begin{bmatrix} 2 & 5 & 8 \\ -1 & 5 & -3 \end{bmatrix} \quad \vec{v} = \begin{bmatrix} 2 \\ -1 \\ -5 \end{bmatrix}$$

Calculation for the first row of A :

$$2 \cdot 2 + 5(-1) + 8(-5) = 4 - 5 - 40 = -41$$

Calculation for the second row of A :

$$(-1)(2) + (5)(-1) + (-3)(-5) = -2 - 5 + 15 = 8$$

The result vector is $\begin{bmatrix} -41 \\ 8 \end{bmatrix}$.

$$w_1 \vec{v} + w_2 \vec{v}$$

Another view of matrix-vector multiplication

- Another way of thinking about the product $A\vec{v}$ is that it is a linear combination of the columns of A , using the weights in $\vec{v}$.
- Example: What is $A\vec{v}$?

weights
scalars

$$A = \begin{bmatrix} 2 & 5 & 8 \\ -1 & 5 & -3 \end{bmatrix} \quad \vec{v} = \begin{bmatrix} 2 \\ -1 \\ -5 \end{bmatrix}$$

$A\vec{v} = 2 \begin{bmatrix} 2 \\ -1 \end{bmatrix} + (-1) \begin{bmatrix} 5 \\ 5 \end{bmatrix} + (-5) \begin{bmatrix} 8 \\ -3 \end{bmatrix} = \begin{bmatrix} -4 \\ 1 \end{bmatrix}$

vectors are cols of A

linear combination of cols of A

Matrix-vector products create linear combinations of columns!

- **Key idea:** It'll be very useful to think of the matrix-vector product $A\vec{v}$ as a linear combination of the columns of A , using the weights in $\vec{v}$.

$$A = \begin{bmatrix} a_{11} & a_{12} & \dots & a_{1d} \\ a_{21} & a_{22} & \dots & a_{2d} \\ \vdots & \vdots & \vdots & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nd} \end{bmatrix} \quad \vec{v} = \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_d \end{bmatrix}$$

$n \times d$ $d \times 1$



$$A\vec{v} = v_1 \begin{bmatrix} a_{11} \\ a_{21} \\ \vdots \\ a_{n1} \end{bmatrix} + v_2 \begin{bmatrix} a_{12} \\ a_{22} \\ \vdots \\ a_{n2} \end{bmatrix} + \dots + v_d \begin{bmatrix} a_{1d} \\ a_{2d} \\ \vdots \\ a_{nd} \end{bmatrix}$$

vector in $\mathbb{R}^n$

~~D~~ come back

Spans and projections, revisited

Moving to multiple dimensions

- Let's now consider three vectors, $\vec{y}$, $\vec{x}^{(1)}$, and $\vec{x}^{(2)}$, all in $\mathbb{R}^n$.

- Question:** What vector in $\text{span}(\vec{x}^{(1)}, \vec{x}^{(2)})$ is closest to $\vec{y}$?

- That is, what values of w_1 and w_2 minimize $\|\vec{e}\| = \|\vec{y} - w_1\vec{x}^{(1)} - w_2\vec{x}^{(2)}\|$?

scalars

w_1 & w_2 such that

$$\vec{x}^{(1)} \cdot \vec{e} = 0$$

$$\vec{x}^{(2)} \cdot \vec{e} = 0$$

$\text{span}(\vec{x}^{(1)}, \vec{x}^{(2)})$

$$w_1\vec{x}^{(1)} + w_2\vec{x}^{(2)} - \vec{y} = \vec{e}$$

Matrix-vector products create linear combinations of columns!

$$\vec{x}^{(1)} = \begin{bmatrix} 2 \\ 5 \\ 3 \end{bmatrix}$$

$$\vec{x}^{(2)} = \begin{bmatrix} -1 \\ 0 \\ 4 \end{bmatrix}$$

$$\vec{y} = \begin{bmatrix} 1 \\ 3 \\ 9 \end{bmatrix}$$

scalars

$$\vec{w} = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}$$

$w_1 \vec{x}^{(1)} + w_2 \vec{x}^{(2)}$

- Combining $\vec{x}^{(1)}$ and $\vec{x}^{(2)}$ into a single matrix gives:

$$X = \begin{bmatrix} | & | \\ \vec{x}^{(1)} & \vec{x}^{(2)} \\ | & | \end{bmatrix} = \begin{bmatrix} 2 & -1 \\ 5 & 0 \\ 3 & 4 \end{bmatrix}$$

Set of vectors

$X\vec{w} = w_1 \begin{bmatrix} 2 \\ 5 \\ 3 \end{bmatrix} + w_2 \begin{bmatrix} -1 \\ 0 \\ 4 \end{bmatrix}$

- Then, if $\vec{w} = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}$, linear combinations of $\vec{x}^{(1)}$ and $\vec{x}^{(2)}$ can be written as $X\vec{w}$.
- The **span of the columns of X** , or $\text{span}(X)$, consists of all vectors that can be written in the form $X\vec{w}$.

Minimizing projection error in multiple dimensions

$$X = \begin{bmatrix} | & | \\ \vec{x}^{(1)} & \vec{x}^{(2)} \\ | & | \end{bmatrix} = \begin{bmatrix} 2 & -1 \\ 5 & 0 \\ 3 & 4 \end{bmatrix} \quad \vec{y} = \begin{bmatrix} 1 \\ 3 \\ 9 \end{bmatrix}$$

- **Goal:** Find the vector $\vec{w} = [w_1 \ w_2]^T$ such that $\|\vec{e}\| = \|\vec{y} - X\vec{w}\|$ is minimized.
- As we've seen, $\vec{w}$ must be such that:

$$\left\{ \begin{array}{l} \vec{x}^{(1)} \cdot \left(\vec{y} - w_1 \vec{x}^{(1)} - w_2 \vec{x}^{(2)} \right) = 0 \\ \vec{x}^{(2)} \cdot \underbrace{\left(\vec{y} - w_1 \vec{x}^{(1)} - w_2 \vec{x}^{(2)} \right)}_{\vec{e}} = 0 \end{array} \right.$$

- How can we use our knowledge of matrices to rewrite this system of equations as a single equation?

Simplifying the system of equations, using matrices

$$X = \begin{bmatrix} | & | \\ \vec{x}^{(1)} & \vec{x}^{(2)} \\ | & | \end{bmatrix} = \begin{bmatrix} 2 & -1 \\ 5 & 0 \\ 3 & 4 \end{bmatrix} \quad \vec{y} = \begin{bmatrix} 1 \\ 3 \\ 9 \end{bmatrix}$$

$$\begin{aligned} \vec{x}^{(1)} \cdot (\vec{y} - w_1 \vec{x}^{(1)} - w_2 \vec{x}^{(2)}) &= 0 \\ \vec{x}^{(2)} \cdot (\vec{y} - w_1 \vec{x}^{(1)} - w_2 \vec{x}^{(2)}) &= 0 \end{aligned}$$

$\underbrace{\hspace{10em}}_{\vec{e}}$

$X \vec{w}$

$\Rightarrow \begin{cases} \vec{x}^{(1)} \cdot (\vec{y} - X \vec{w}) = 0 \\ \vec{x}^{(2)} \cdot (\vec{y} - X \vec{w}) = 0 \end{cases}$

Simplifying the system of equations, using matrices

$$X = \begin{bmatrix} | & | \\ \vec{x}^{(1)} & \vec{x}^{(2)} \\ | & | \end{bmatrix} = \begin{bmatrix} 2 & -1 \\ 5 & 0 \\ 3 & 4 \end{bmatrix} \quad \vec{y} = \begin{bmatrix} 1 \\ 3 \\ 9 \end{bmatrix}$$

1. $w_1 \vec{x}^{(1)} + w_2 \vec{x}^{(2)}$ can be written as $X\vec{w}$, so $\vec{e} = \vec{y} - X\vec{w}$.
2. The condition that $\vec{e}$ must be orthogonal to each column of X is equivalent to condition that $X^T \vec{e} = 0$.

$$\begin{bmatrix} \vec{x}^{(1)} \\ \vec{x}^{(2)} \end{bmatrix} \cdot (\vec{y} - X \vec{w}) = 0$$

$$\vec{x}^{(2)} \cdot (\vec{y} - X \vec{w}) = 0$$

Goal: combine
into 1 eq

$$X^T (\vec{y} - X \vec{w}) = 0$$

↑ rows X^T are cols of X
rewrite the equations

1 eq

$$X = \begin{bmatrix} \vec{x}^{(1)} & \vec{x}^{(2)} \\ | & | \end{bmatrix}^{3 \times 2}$$

$$\vec{e} = 0$$

$$X^T = \begin{bmatrix} \vec{x}^{(1)} \\ \vec{x}^{(2)} \end{bmatrix}^{2 \times 3}$$

$$= \begin{bmatrix} 2 & 5 & 3 \\ -1 & 0 & 4 \end{bmatrix}$$

$$\vec{a} \cdot \vec{b} = a^T b =$$

The normal equations

$$X = \begin{bmatrix} | & | \\ \vec{x}^{(1)} & \vec{x}^{(2)} \\ | & | \end{bmatrix} = \begin{bmatrix} 2 & -1 \\ 5 & 0 \\ 3 & 4 \end{bmatrix} \quad \vec{y} = \begin{bmatrix} 1 \\ 3 \\ 9 \end{bmatrix}$$

- **Goal:** Find the vector $\vec{w} = [w_1 \ w_2]^T$ such that $\|\vec{e}\| = \|\vec{y} - X\vec{w}\|$ is minimized.
- We now know that it is the vector $\vec{w}^*$ such that:

$$X^T \vec{e} = 0$$

$$\rightarrow X^T (\vec{y} - X\vec{w}^*) = 0$$

$$\rightarrow X^T \vec{y} - X^T X \vec{w}^* = 0$$

$$\Rightarrow X^T X \vec{w}^* = X^T \vec{y}$$

- The last statement is referred to as the **normal equations**.

normal orthogonal

same idea as
 $\frac{2w}{2} = 0 \Rightarrow w = 0$
get rid of $X^T X$
by taking the inverse?
 $(X^T X)^{-1}$
system of equations

The general solution to the normal equations

$$X \in \mathbb{R}^{n \times d} \quad \vec{y} \in \mathbb{R}^n$$

$$\begin{matrix} 2x + 4y \\ x + 2y \end{matrix} \rightarrow x2$$

- **Goal, in general:** Find the vector $\vec{w} \in \mathbb{R}^d$ such that $\|\vec{e}\| = \|\vec{y} - X\vec{w}\|$ is minimized.
- We now know that it is the vector $\vec{w}^*$ such that:

$$\begin{aligned} X^T \vec{e} &= 0 \\ \implies X^T X \vec{w}^* &= X^T \vec{y} \end{aligned}$$

- Assuming $X^T X$ is invertible, this is the vector:

$$\vec{w}^* = (X^T X)^{-1} X^T \vec{y}$$

no column can
be expressed as
a combination of
the others

- This is a big assumption, because it requires $X^T X$ to be **full rank**.
- If $X^T X$ is not full rank, then there are infinitely many solutions to the normal equations, $X^T X \vec{w}^* = X^T \vec{y}$.

What does it mean?

- **Original question:** What vector in $\text{span}(\vec{x}^{(1)}, \vec{x}^{(2)})$ is closest to $\vec{y}$?
- **Final answer:** It is the vector $X\vec{w}^*$, where:

$$\vec{w}^* = (X^T X)^{-1} X^T \vec{y}$$

- Revisiting our example:

$$X = \begin{bmatrix} | & | \\ \vec{x}^{(1)} & \vec{x}^{(2)} \\ | & | \end{bmatrix} = \begin{bmatrix} 2 & -1 \\ 5 & 0 \\ 3 & 4 \end{bmatrix} \quad \vec{y} = \begin{bmatrix} 1 \\ 3 \\ 9 \end{bmatrix}$$

- Using a computer gives us $\vec{w}^* = (X^T X)^{-1} X^T \vec{y} \approx \begin{bmatrix} 0.7289 \\ 1.6300 \end{bmatrix}$.
- So, the vector in $\text{span}(\vec{x}^{(1)}, \vec{x}^{(2)})$ closest to $\vec{y}$ is $0.7289\vec{x}^{(1)} + 1.6300\vec{x}^{(2)}$.

An optimization problem, solved

- We just used linear algebra to solve an **optimization problem**.
- Specifically, the function we minimized is:

$$\text{error}(\vec{w}) = \|\vec{y} - X\vec{w}\|$$

- This is a function whose input is a vector, $\vec{w}$, and whose output is a scalar!
- The input, $\vec{w}^*$, to $\text{error}(\vec{w})$ that minimizes it is:

$$\vec{w}^* = (X^T X)^{-1} X^T \vec{y}$$

- We're going to use this frequently!

Announcements

Announcements

- Homework 3 is due on **Friday, July 10th**, 11:59pm PT.
 - Try to finish it relatively early. Late submissions receive 0 credit, and we had a couple on HW1. Please don't make us give you a 0.
- Homework 1 scores will be available on Gradescope soon (hopefully), we will email you when this is the case.
 - Keep an eye out because the regrade request window is short.

Exam Announcements

- Exam Proctoring Instructions have been sent out. Read them in full, carefully.
 - It is **your** responsibility to read them and ensure you have a proper proctoring set up. We will not be waiting to start the exam if you are late or are not set up; your exam score will just be canceled.
 - We are being very strict with proctoring for exam security, and any student caught without the correct proctoring set is eligible to having their exam score canceled.
- Many of you are in non-PT time zones. Please plan to take the exam during the expected exam time slot: Wednesday, July 15th, 11:00AM-1:50PM, PT.
 - We will **not** be accommodating different exam times, unless you have OSD accommodations specifically for that.
 - We will be confirming the length of the exam today or tomorrow.

90mins 11:00 - 12:48 PM

Exam Announcements

- If you intend to use your exam accommodations on this upcoming midterm, please make sure OSD sends your AFA letter to us **by this Friday, July 10th**.
 - It's hard to change/reschedule your exam last minute, so giving us at least the weekend to rearrange and make sure we can accommodate you is really helpful.
 - If you have done this, you should be hearing from us soon if you haven't already to confirm your exam accommodations.
- If you have any questions about the exam or proctoring, please post on Piazza.

Jordan:
mention gap matching far

12:36 pm

Break

Lecture 8

Regression and Linear Algebra

DSC 40A, Summer 2024

Announcements

- Homework 3 is due on **Saturday, April 27th**.
 - We moved some office hours around – we now have some on Saturday!
- Homework 1 scores are available on Gradescope.
 - Regrade requests are due on Sunday.
- Groupwork 4 is on Monday. Remember to submit groupworks as a group – you won't get any credit if you work alone!
- The Midterm Exam is on **Tuesday, May 7th in class**.
 - We will have a review session on **Friday, May 3rd from 2-5PM** where we'll go over old homework and exam problems.
 - We will be posting many past exams this weekend!

Agenda

- Overview: Spans and projections.
- Regression and linear algebra.
- Multiple linear regression.

Question 🤔

Answer in chat

Remember, you can always ask questions in Q&A"


Overview: Spans and projections

Projecting onto the span of a single vector

- **Question:** What vector in $\text{span}(\vec{x})$ is closest to $\vec{y}$?
- The answer is the vector $w\vec{x}$, where the w is chosen to minimize the **length** of the error vector:

$$\|\vec{e}\| = \|\vec{y} - w\vec{x}\|$$


- **Key idea:** To minimize the length of the error vector, choose w so that the error vector is **orthogonal** to $\vec{x}$.



Projecting onto the span of a single vector

- **Question:** What vector in $\text{span}(\vec{x})$ is closest to $\vec{y}$?
- **Answer:** It is the vector $w^* \vec{x}$, where:

$$w^* = \frac{\vec{x} \cdot \vec{y}}{\vec{x} \cdot \vec{x}}$$




Projecting onto the span of multiple vectors

- **Question:** What vector in $\text{span}(\vec{x}^{(1)}, \vec{x}^{(2)})$ is closest to $\vec{y}$?

- The answer is the vector $w_1\vec{x}^{(1)} + w_2\vec{x}^{(2)}$, where w_1 and w_2 are chosen to minimize the **length** of the error vector:

$$\|\vec{e}\| = \|\vec{y} - w_1\vec{x}^{(1)} - w_2\vec{x}^{(2)}\|$$

- **Key idea:** To minimize the length of the error vector, choose w_1 and w_2 so that the error vector is **orthogonal** to **both** $\vec{x}^{(1)}$ and $\vec{x}^{(2)}$.



If $\vec{x}^{(1)}$ and $\vec{x}^{(2)}$ are **linearly independent**, they span a **plane**.

Matrix-vector products create linear combinations of columns!

- **Question:** What vector in $\text{span}(\vec{x}^{(1)}, \vec{x}^{(2)})$ is closest to $\vec{y}$?
- To help, we can create a **matrix**, X , by stacking $\vec{x}^{(1)}$ and $\vec{x}^{(2)}$ next to each other:

$$X = \begin{bmatrix} | & | \\ \vec{x}^{(1)} & \vec{x}^{(2)} \\ | & | \end{bmatrix} = \begin{bmatrix} 2 & -1 \\ 5 & 0 \\ 3 & 4 \end{bmatrix} \quad \vec{y} = \begin{bmatrix} 1 \\ 3 \\ 9 \end{bmatrix}$$

- Then, instead of writing vectors in $\text{span}(\vec{x}^{(1)}, \vec{x}^{(2)})$ as $w_1\vec{x}^{(1)} + w_2\vec{x}^{(2)}$, we can say:

$$X\vec{w} \quad \text{where } \vec{w} = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}$$

- **Key idea:** Find $\vec{w}$ such that the error vector, $\vec{e} = \vec{y} - X\vec{w}$, is **orthogonal** to every column of X .

Constructing an orthogonal error vector

- **Key idea:** Find $\vec{w} \in \mathbb{R}^d$ such that the error vector, $\vec{e} = \vec{y} - X\vec{w}$, is **orthogonal** to the columns of X .

- Why? Because this will make the error vector as short as possible.

- The $\vec{w}^*$ that accomplishes this satisfies:

vector

$$X^T \vec{e} = 0 \quad \text{0 vector} \quad \begin{bmatrix} 0 \\ \vdots \\ 0 \end{bmatrix}$$

$[\vec{y} - X\vec{w}^*] = 0$

- Why? Because $X^T \vec{e}$ contains the **dot products** of each column in X with $\vec{e}$. If these are all 0, then $\vec{e}$ is **orthogonal to every column of X** !

$$X = \begin{bmatrix} 1 & x^{(1)} \\ x_{(1)} & x^{(2)} \\ 1 & 1 \end{bmatrix}_{3 \times 2}$$

$$X^T \vec{e} = \begin{bmatrix} -\vec{x}^{(1)T} \\ -\vec{x}^{(2)T} \end{bmatrix} \vec{e} = \begin{bmatrix} \vec{x}^{(1)T} \vec{e} \\ \vec{x}^{(2)T} \vec{e} \end{bmatrix} = 0$$

$A \vec{v}$ $\rightarrow$ dot product of $\vec{v}$ w/ rows of A

The normal equations

- **Key idea:** Find $\vec{w} \in \mathbb{R}^d$ such that the error vector, $\vec{e} = \vec{y} - X\vec{w}$, is **orthogonal** to the columns of X .
- The $\vec{w}^*$ that accomplishes this satisfies:

$$\begin{aligned}X^T \vec{e} &= 0 \\X^T (\vec{y} - X\vec{w}^*) &= 0 \\X^T \vec{y} - X^T X \vec{w}^* &= 0 \\ \implies X^T X \vec{w}^* &= X^T \vec{y}\end{aligned}$$

multiply by inverse

$$\vec{w}^* = (X^T X)^{-1} X^T \vec{y}$$

- This is a big assumption, because it requires $X^T X$ to be **full rank**. *linearly indep cols*
- If $X^T X$ is not full rank, then there are infinitely many solutions to the normal equations, $X^T X \vec{w}^* = X^T \vec{y}$.

Orthogonal system of eq

- The last statement is referred to as the **normal equations**.

What does it mean?

- **Original question:** What vector in $\text{span}(\vec{x}^{(1)}, \vec{x}^{(2)})$ is closest to $\vec{y}$?
- **Final answer:** Assuming $X^T X$ is invertible, it is the vector $X\vec{w}^*$, where:

$$\vec{w}^* = (X^T X)^{-1} X^T \vec{y}$$

- Revisiting our example:

$$X = \begin{bmatrix} | & | \\ \vec{x}^{(1)} & \vec{x}^{(2)} \\ | & | \end{bmatrix} = \begin{bmatrix} 2 & -1 \\ 5 & 0 \\ 3 & 4 \end{bmatrix} \quad \vec{y} = \begin{bmatrix} 1 \\ 3 \\ 9 \end{bmatrix}$$

- Using a computer gives us $\vec{w}^* = (X^T X)^{-1} X^T \vec{y} \approx \begin{bmatrix} 0.7289 \\ 1.6300 \end{bmatrix}$
- So, the vector in $\text{span}(\vec{x}^{(1)}, \vec{x}^{(2)})$ closest to $\vec{y}$ is $0.7289\vec{x}^{(1)} + 1.6300\vec{x}^{(2)}$.

Best
Slope

An optimization problem, solved

- We just used linear algebra to solve an **optimization problem**.

- Specifically, the function we minimized is:

$$\text{error}(\vec{w}) = \|\vec{y} - X\vec{w}\|$$

MSE/MAE
how close can I make it?

- This is a function whose input is a vector, $\vec{w}$, and whose output is a scalar!
- The input, $\vec{w}^*$, to $\text{error}(\vec{w})$ that minimizes it is one that satisfies the **normal equations**:

$$X^T X \vec{w}^* = X^T \vec{y}$$

If $X^T X$ is invertible, then the unique solution is:

$$\vec{w}^* = (X^T X)^{-1} X^T \vec{y}$$


- We're going to use this frequently!

Regression and linear algebra

Wait... why do we need linear algebra?

- Soon, we'll want to make predictions using more than one feature.
 - Example: Predicting commute times using departure hour and temperature.
- Thinking about linear regression in terms of **matrices and vectors** will allow us to find hypothesis functions that: *more complicated*
 - Use multiple features (input variables).
 - Are non-linear in the features, e.g. $H(x) = w_0 + w_1x + w_2x^2$.
- Let's see if we can put what we've just learned to use.

Simple linear regression, revisited



- **Model:** $H(x) = w_0 + w_1 x$.
Intercept (pointing to w_0) and *Slope* (pointing to w_1)
- **Loss function:** $(y_i - H(x_i))^2$.
LSQ (Least Squares)
- To find w_0^* and w_1^* , we minimized empirical risk, i.e. average loss:
Best (pointing to w_0^* and w_1^*)

$$R_{\text{sq}}(H) = \frac{1}{n} \sum_{i=1}^n (y_i - H(x_i))^2$$

Best (pointing to the loss function)

- **Observation:** $R_{\text{sq}}(w_0, w_1)$ kind of looks like the formula for the norm of a vector,

$$\|\vec{v}\| = \sqrt{v_1^2 + v_2^2 + \dots + v_n^2}$$

Norm (pointing to the formula)

Vocab

Regression and linear algebra

Let's define a few new terms:

$$\begin{bmatrix} 42 \text{ mths} \\ \vdots \end{bmatrix}_{n \times 1}$$

- The **observation vector** is the vector $\vec{y} \in \mathbb{R}^n$. This is the vector of observed "actual values". y_i

- The **hypothesis vector** is the vector $\vec{h} \in \mathbb{R}^n$ with components $H(x_i)$. This is the vector of predicted values.

← function


- The **error vector** is the vector $\vec{e} \in \mathbb{R}^n$ with components:

$$e_i = y_i - H(x_i)$$
$$\vec{e} = \vec{y} - \vec{h}$$

$$\begin{bmatrix} 42 \text{ mths} \\ \vdots \end{bmatrix}_{n \times 1}$$

Example

Consider $H(x) = 2 + \frac{1}{2}x$.



Observation vector

$$\vec{y} = \begin{bmatrix} 4 \\ 5 \\ 2 \end{bmatrix}$$

$$\vec{h} = \begin{bmatrix} 2 + \frac{1}{2}(1) \\ 2 + \frac{1}{2}(5) \\ 2 + \frac{1}{2}(4) \end{bmatrix} = \begin{bmatrix} \frac{5}{2} \\ \frac{9}{2} \\ 4 \end{bmatrix}$$

$$\vec{e} = \vec{y} - \vec{h} = \begin{bmatrix} 4 - \frac{5}{2} \\ 5 - \frac{9}{2} \\ 2 - 4 \end{bmatrix} = \begin{bmatrix} \frac{3}{2} \\ \frac{1}{2} \\ -2 \end{bmatrix}$$

error vector

$$R_{sq}(H) = \frac{1}{n} \sum_{i=1}^n (y_i - H(x_i))^2 = \frac{1}{3} \left[\left(\frac{3}{2}\right)^2 + \left(\frac{1}{2}\right)^2 + (-2)^2 \right]$$

$$R_{sq}(H) = \frac{1}{3} \|\vec{e}\|^2$$

length of error vector

Regression and linear algebra

Let's define a few new terms:

- The **observation vector** is the vector $\vec{y} \in \mathbb{R}^n$. This is the vector of observed "actual values".
- The **hypothesis vector** is the vector $\vec{h} \in \mathbb{R}^n$ with components $H(x_i)$. This is the vector of predicted values.
- The **error vector** is the vector $\vec{e} \in \mathbb{R}^n$ with components:

$$e_i = y_i - H(x_i)$$

- **Key idea:** We can rewrite the mean squared error of H as:

$$R_{\text{sq}}(H) = \frac{1}{n} \sum_{i=1}^n (y_i - H(x_i))^2 = \frac{1}{n} \|\vec{e}\|^2 = \frac{1}{n} \|\vec{y} - \vec{h}\|^2$$

Handwritten notes: "connection!" (pointing to the error vector definition), "actual" (pointing to $\vec{y}$), "predicted" (pointing to $\vec{h}$)

The hypothesis vector

- The **hypothesis vector** is the vector $\vec{h} \in \mathbb{R}^n$ with components $H(x_i)$. This is the vector of predicted values.
- For the linear hypothesis function $H(x) = w_0 + w_1 x$, the hypothesis vector can be written:

intercept slope

hypothesis vector

$$\vec{h} = \begin{bmatrix} w_0 + w_1 x_1 \\ w_0 + w_1 x_2 \\ \vdots \\ w_0 + w_1 x_n \end{bmatrix} = \begin{bmatrix} \underbrace{1}_{\text{intercept}} & \underbrace{x_1}_{\text{slope}} \\ \underbrace{1} & \underbrace{x_2} \\ \vdots & \vdots \\ \underbrace{1} & \underbrace{x_n} \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \end{bmatrix}$$

"Design matrix"

$1 \times n$ $n \times 2$ 2×1

$w_0 + w_1 x$

Rewriting the mean squared error

- Define the **design matrix** $X \in \mathbb{R}^{n \times 2}$ as:

$$X = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_n \end{bmatrix}$$

Still

$$H(x_i) = w_0 + w_1 x_i$$

↑

- Define the **parameter vector** $\vec{w} \in \mathbb{R}^2$ to be $\vec{w} = \begin{bmatrix} w_0 \\ w_1 \end{bmatrix}$

intercept

slope

- Then $\vec{h} = X\vec{w}$, so the mean squared error becomes:

$$R_{\text{sq}}(\vec{H}) = \frac{1}{n} \|\vec{y} - \vec{h}\|^2 \implies R_{\text{sq}}(\vec{w}) = \frac{1}{n} \|\vec{y} - X\vec{w}\|^2$$

Minimizing mean squared error, again

- To find the optimal model parameters for simple linear regression, w_0^* and w_1^* , we previously minimized:

$$R_{\text{sq}}(w_0, w_1) = \frac{1}{n} \sum_{i=1}^n (y_i - (w_0 + w_1 x_i))^2$$

- Now that we've reframed the simple linear regression problem in terms of linear algebra, we can find w_0^* and w_1^* by finding the $\vec{w}^* = [w_0^* \quad w_1^*]^T$ that minimizes:

$$R_{\text{sq}}(\vec{w}) = \frac{1}{n} \|\vec{y} - X\vec{w}\|^2$$

- Do we already know the $\vec{w}^*$ that minimizes $R_{\text{sq}}(\vec{w})$?

we know how to minimize

An optimization problem we've seen before

best: intercept Slope

$$\frac{1}{n} \sum (y_i - \hat{y}_i)^2$$

- The optimal parameter vector, $\vec{w}^* = [w_0^* \ w_1^*]^T$, is the one that minimizes:

$$R_{\text{sq}}(\vec{w}) = \frac{1}{n} \|\vec{y} - X\vec{w}\|^2$$

close relationship
MSE & $\|\vec{e}\|$

prev
slide

- Previously, we found that $\vec{w}^* = (X^T X)^{-1} X^T \vec{y}$ minimizes the length of the error vector, $\|\vec{e}\| = \|\vec{y} - X\vec{w}\|$

middle of lecture

- $R_{\text{sq}}(\vec{w})$ is closely related to $\|\vec{e}\|$:

not $\&$ $R_{\text{sq}}(\vec{w}) = \frac{1}{n} \|\vec{e}\|^2$
just $\|\vec{e}\|$

- The minimizer of $\|\vec{e}\|$ is the same as the minimizer of $R_{\text{sq}}(\vec{w})$!
- Key idea:** $\vec{w}^* = (X^T X)^{-1} X^T \vec{y}$ also minimizes $R_{\text{sq}}(\vec{w})$!

The optimal parameter vector, $\vec{w}^*$

- To find the optimal model parameters for simple linear regression, w_0^* and w_1^* , we previously minimized $R_{\text{sq}}(w_0, w_1) = \frac{1}{n} \sum_{i=1}^n (y_i - (w_0 + w_1 x_i))^2$.

- We found, using calculus, that:

1,000ws
1000 eas
:

$$w_1^* = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2} = r \frac{\sigma_y}{\sigma_x}$$

$$w_0^* = \bar{y} - w_1^* \bar{x}$$

← optimal intercept

- Another way of finding optimal model parameters for simple linear regression is to find the $\vec{w}^*$ that minimizes $R_{\text{sq}}(\vec{w}) = \frac{1}{n} \|\vec{y} - X\vec{w}\|^2$. *same things written differently*

- MSE*
- The minimizer, if $X^T X$ is invertible, is the vector $\vec{w}^* = (X^T X)^{-1} X^T \vec{y}$.

- These formulas are equivalent!

Roadmap

- To give us a break from math, we'll switch to a notebook, [linked here](#), showing that both formulas – that is, (1) the formulas for w_1^* and w_0^* we found using calculus, and (2) the formula for $\vec{w}^*$ we found using linear algebra – give the same results.
- Then, we'll use our new linear algebraic formulation of regression to incorporate **multiple features** in our prediction process.

Summary: Regression and linear algebra

- Define the **design matrix** $X \in \mathbb{R}^{n \times 2}$, **observation vector** $\vec{y} \in \mathbb{R}^n$, and parameter vector $\vec{w} \in \mathbb{R}^2$ as:

$$X = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_n \end{bmatrix} \quad \vec{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} \quad \vec{w} = \begin{bmatrix} w_0 \\ w_1 \end{bmatrix}$$

- How do we make the hypothesis vector, $\vec{h} = X\vec{w}$, as close to $\vec{y}$ as possible? Use the parameter vector $\vec{w}^*$:

$$\vec{w}^* = (X^T X)^{-1} X^T \vec{y}$$

- We chose $\vec{w}^*$ so that $\vec{h}^* = X\vec{w}^*$ is the projection of $\vec{y}$ onto the span of the columns of the design matrix, X .

Multiple linear regression

	departure_hour	day_of_month	minutes
0	10.816667	15	68.0
1	7.750000	16	94.0
2	8.450000	22	63.0
3	7.133333	23	100.0
4	9.150000	30	69.0
...	...	...	...

So far, we've fit **simple** linear regression models, which use only **one** feature (`'departure_hour'`) for making predictions.

Incorporating multiple features

- In the context of the commute times dataset, the simple linear regression model we fit was of the form:

$$\begin{aligned}\text{pred. commute} &= H(\text{departure hour}) \\ &= w_0 + w_1 \cdot \text{departure hour}\end{aligned}$$

- Now, we'll try and fit a multiple linear regression model of the form:

$$\begin{aligned}\text{pred. commute} &= H(\text{departure hour}) \\ &= w_0 + w_1 \cdot \text{departure hour} + w_2 \cdot \text{day of month}\end{aligned}$$

- Linear regression with **multiple** features is called **multiple linear regression**.
- How do we find w_0^* , w_1^* , and w_2^* ?

Geometric interpretation

- The hypothesis function:

$$H(\text{departure hour}) = w_0 + w_1 \cdot \text{departure hour}$$

looks like a **line** in 2D.


- **Questions:**

- How many dimensions do we need to graph the hypothesis function:

$$H(\text{departure hour}) = w_0 + w_1 \cdot \text{departure hour} + w_2 \cdot \text{day of month}$$

- What is the shape of the hypothesis function?

Commute Time vs. Departure Hour and Day of Month



Our new hypothesis function is a **plane** in 3D!

The setup

- Suppose we have the following dataset.

	departure_hour	day_of_month	minutes
row			
1	8.45	22	63.0
2	8.90	28	89.0
3	8.72	18	89.0

- We can represent each day with a **feature vector**, $\vec{x}$:

The hypothesis vector

- When our hypothesis function is of the form:

$$H(\text{departure hour}) = w_0 + w_1 \cdot \text{departure hour} + w_2 \cdot \text{day of month}$$

the hypothesis vector $\vec{h} \in \mathbb{R}^n$ can be written as:

$$\vec{h} = \begin{bmatrix} H(\text{departure hour}_1, \text{day}_1) \\ H(\text{departure hour}_2, \text{day}_2) \\ \dots \\ H(\text{departure hour}_n, \text{day}_n) \end{bmatrix} = \begin{bmatrix} 1 & \text{departure hour}_1 & \text{day}_1 \\ 1 & \text{departure hour}_2 & \text{day}_2 \\ \dots & \dots & \dots \\ 1 & \text{departure hour}_n & \text{day}_n \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \\ w_2 \end{bmatrix}$$

Finding the optimal parameters

- To find the optimal parameter vector, $\vec{w}^*$, we can use the **design matrix** $X \in \mathbb{R}^{n \times 3}$ and **observation vector** $\vec{y} \in \mathbb{R}^n$:

$$X = \begin{bmatrix} 1 & \text{departure hour}_1 & \text{day}_1 \\ 1 & \text{departure hour}_2 & \text{day}_2 \\ \dots & \dots & \dots \\ 1 & \text{departure hour}_n & \text{day}_n \end{bmatrix} \quad \vec{y} = \begin{bmatrix} \text{commute time}_1 \\ \text{commute time}_2 \\ \vdots \\ \text{commute time}_n \end{bmatrix}$$

- Then, all we need to do is solve the **normal equations**:

$$X^T X \vec{w}^* = X^T \vec{y}$$

If $X^T X$ is invertible, we know the solution is:

$$\vec{w}^* = (X^T X)^{-1} X^T \vec{y}$$

Roadmap

- To wrap up today's lecture, we'll find the optimal parameter vector $\vec{w}^*$ for our new two-feature model in code. We'll switch back to our notebook, [linked here](#).
- Next class, we'll present a more general framing of the multiple linear regression model, that uses d features instead of just two.
- We'll also look at how we can **engineer** new features using existing features.
 - e.g. How can we fit a hypothesis function of the form
$$H(x) = w_0 + w_1x + w_2x^2?$$